

Introduction to MongoDB



Raviteja Dodda
Co-Founder, Pipal Tech Ventures

raviteja@pipaltechventures.com
Twitter, Facebook - @raviteja2007

April 1st, 2012

The Great Divide



MongoDB - Sweet Spot: **Easy**, **Flexible** and **Scalable**



What is MongoDB ?

- Scalable High-Performance Open-source, Document-orientated database.
- Built for Speed
- Rich Document based queries for **Easy readability**.
- Full Index Support for **High Performance**.
- Replication and Failover for **High Availability**.
- Auto Sharding for **Easy Scalability**.
- Map / Reduce for **Aggregation**.



Why use MongoDB?

- SQL was invented in the 70's to store **data**.
- MongoDB stores **documents (or) objects**.
- Now-a-days, everyone works with **objects** (Python/Ruby/Java/etc.)
- And we need Databases to persist our **objects**.
Then why not store **objects** directly ?
- Embedded documents and arrays reduce need for joins. **No Joins** and No-multi document **transactions**.



What is MongoDB great for?

- RDBMS replacement for **Web Applications**.
- **Semi-structured** Content Management.
- **Real-time** Analytics & High-Speed Logging.
- Caching and **High Scalability**

Web 2.0, Media, SAAS, Gaming

HealthCare, Finance, Telecom, Government



Not great for?

- Highly **Transactional** Applications.
- Problems requiring **SQL**.

Some Companies using MongoDB in Production



Let's Dive in !



When I say
Database



Think
Database

- Made up of Multiple **Collections**.
- Created **on-the-fly** when referenced for the first time.

When I say
Collection



Think
Table

- Schema-less, and contains **Documents**.
- **Indexable** by one/more keys.
- Created **on-the-fly** when referenced for the first time.
- **Capped Collections**: Fixed size, older records get dropped after reaching the limit.

When I say
Document



Think
Record/Row

- Stored in a **Collection**.
- Can have **_id** key – works like Primary keys in MySQL.
- Supported Relationships – **Embedded (or) References**.
- Document storage in **BSON** (Binary form of JSON).

Understanding the Document Model.

```
var p = {  
  '_id': '3432',  
  'author': DBRef('User', 2),  
  'title': 'Introduction to MongoDB',  
  'body': 'MongoDB is an open sources.. ',  
  'timestamp': Date('01-04-12'),  
  'tags': ['MongoDB', 'NoSQL'],  
  'comments': [{ 'author': DBRef('User', 4),  
                  'date': Date('02-04-12'),  
                  'text': 'Did you see.. ',  
                  'upvotes': 7, ... }  
]  
}  
> db.posts.save(p);
```



Secondary Indexes

Create Index on any field in the document

// 1 means ascending, -1 means descending

```
> db.posts.ensureIndex({'author': 1});
```

//Index Nested Documents

```
> db.posts.ensureIndex('comments.author': 1);
```

// Index on tags

```
> db.posts.ensureIndex({'tags': 1});
```

// Geo-spatial Index

```
> db.posts.ensureIndex({'author.location': '2d'});
```



What about Queries? So Simple

// find posts which has 'MongoDB' tag.

```
> db.posts.find({tags: 'MongoDB'});
```

// find posts by author's comments.

```
> db.posts.find({'comments.author':  
DBRef('User',2)}).count();
```

// find posts written after 31st March.

```
> db.posts.find({'timestamp': {'gte': Date('31-03-12')}});
```

// find posts written by authors around [22, 42]

```
> db.posts.find({'author.location': {'near':[22, 42]}});
```

\$gt, \$lt, \$gte, \$lte, \$ne, \$all, \$in, \$nin, count, limit, skip, group, etc...



What about Updates? **Atomic Operations** makes it simple

```
db.posts.update({_id: '3432'},  
{'title': 'Introduction to MongoDB (updated)',  
'text': 'Updated text',  
${addToSet: {'tags': 'webinar'}});
```

\$set, \$unset

\$push, \$pull, \$pop, \$addToSet

\$inc, \$decr, many more...

Where are my joins and transactions? !!!



Some Cool features

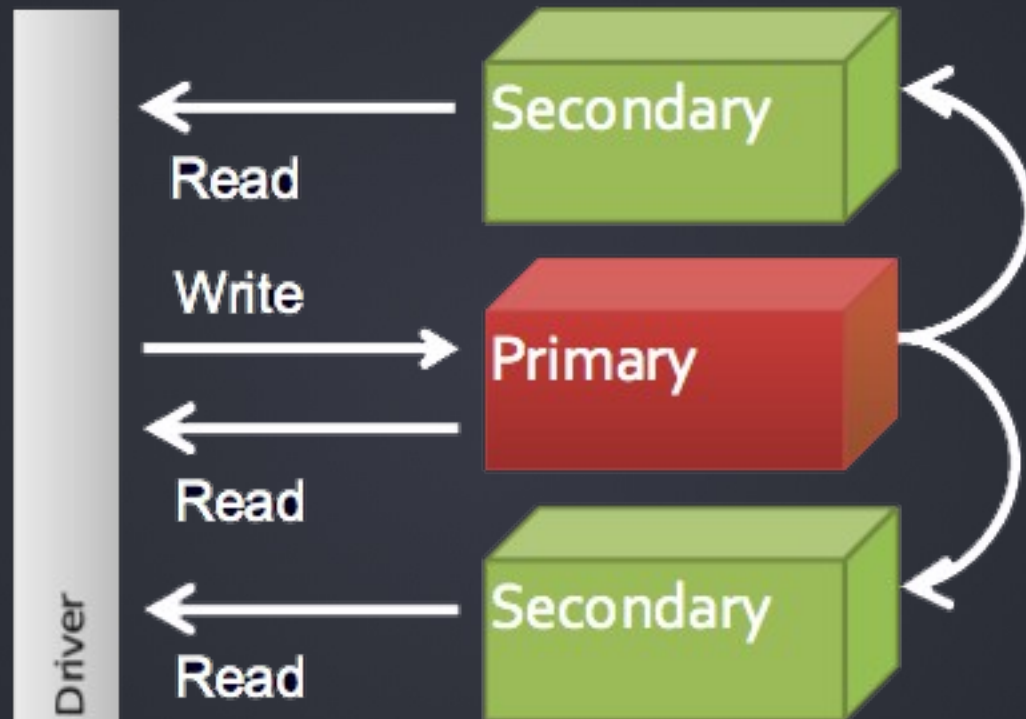
- Geo-spatial Indexes for Geo-spatial queries.
\$near, \$within_distance, Bound queries (circle, box)
- GridFS
Stores Large Binary Files.
- Map/Reduce
GROUP BY in SQL, map/reduce in MongoDB.



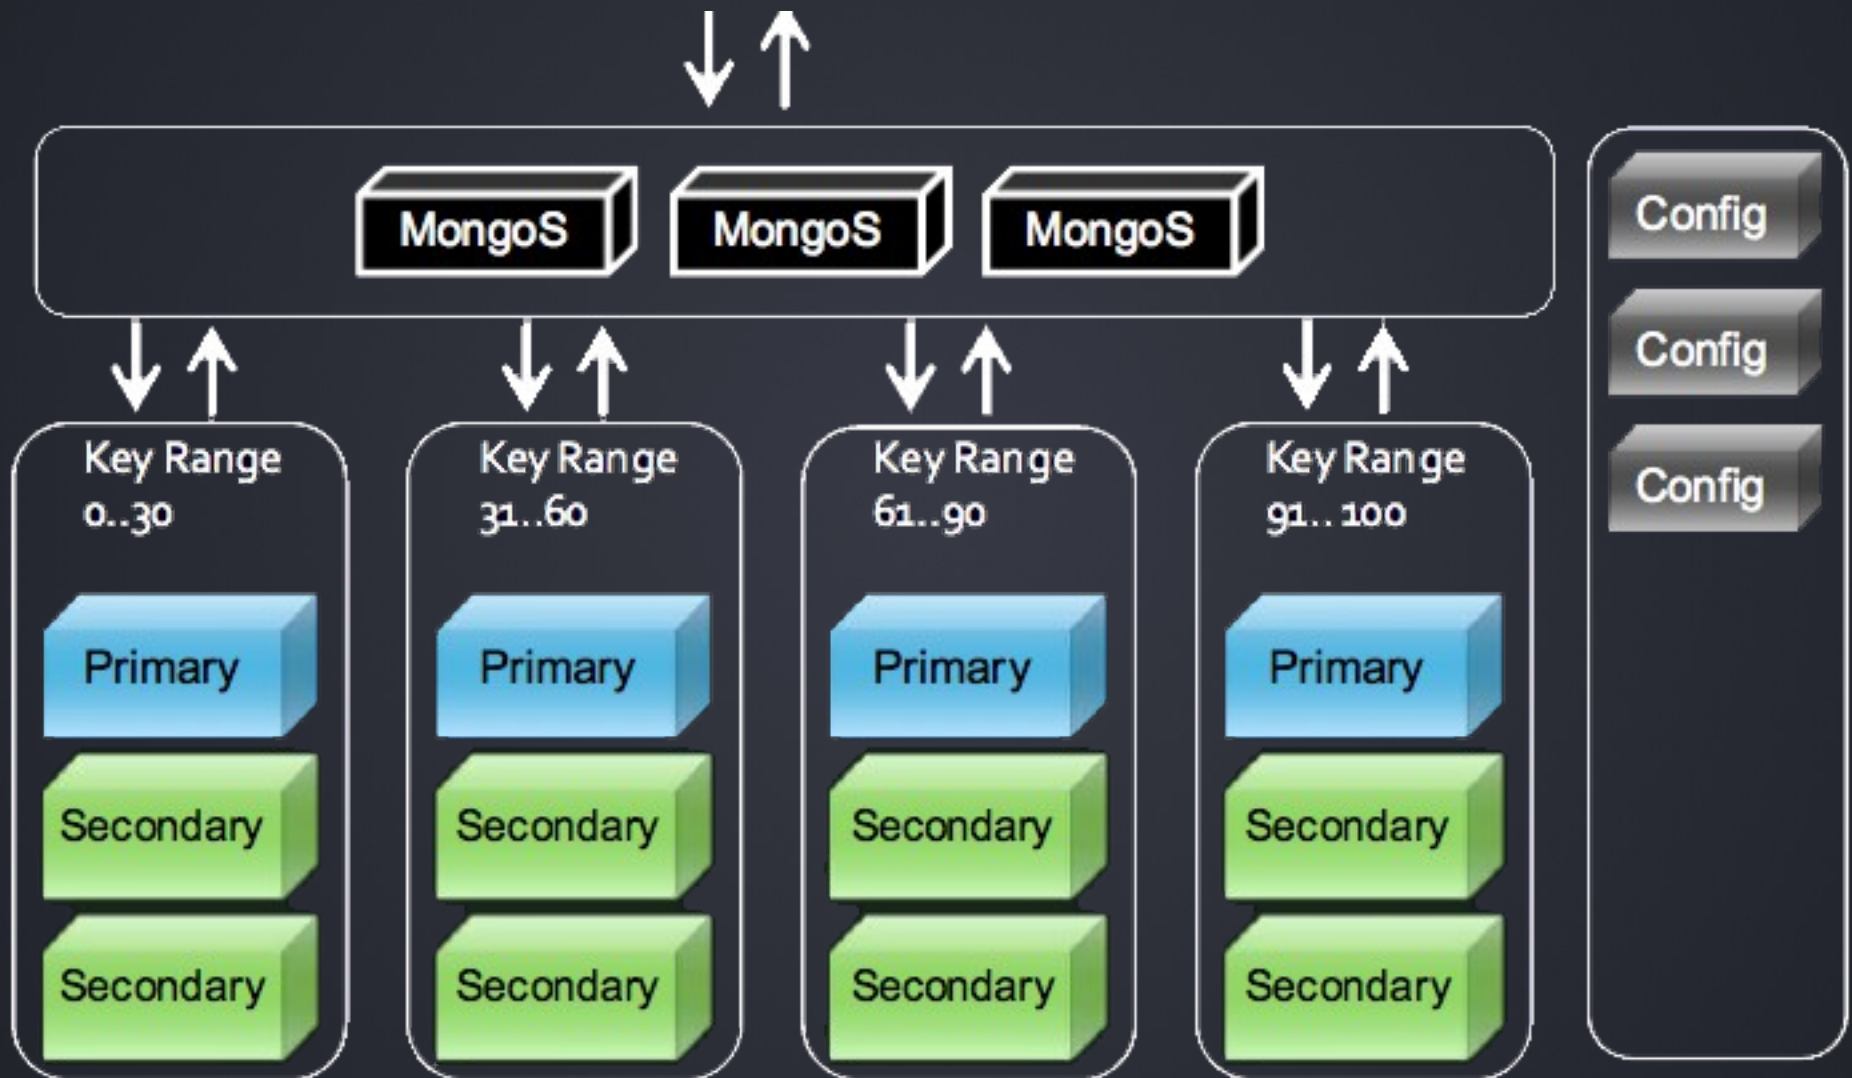
Deployment & Scaling



Replica Sets



Sharding



How do we use MongoDB at Pipal



PyMongo based ORM – A Separate Topic !

Btw, Pipal is hiring at Bangalore

Backend Engineers, Frontend Engineers, System-Administrators

Send us your resume at careers@pipaltechventures.com

Questions?

Next Steps: <http://mongodb.org>,
Twitter: @mongodb

Thank You ☺

Stay Hungry, Stay Foolish !!!

- Steve Jobs

